

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.linear_model import LogisticRegression
```

```
In [2]: df=pd.read_csv("C4_framingham - C4_framingham.csv")
df
```

```
Out[2]:
```

|      | male | age | education | currentSmoker | cigsPerDay | BPMeds | prevalentStroke | prevalentHyp |
|------|------|-----|-----------|---------------|------------|--------|-----------------|--------------|
| 0    | 1    | 39  | 4.0       | 0             | 0.0        | 0.0    | 0               | 0            |
| 1    | 0    | 46  | 2.0       | 0             | 0.0        | 0.0    | 0               | 0            |
| 2    | 1    | 48  | 1.0       | 1             | 20.0       | 0.0    | 0               | 0            |
| 3    | 0    | 61  | 3.0       | 1             | 30.0       | 0.0    | 0               | 1            |
| 4    | 0    | 46  | 3.0       | 1             | 23.0       | 0.0    | 0               | 0            |
| ...  | ...  | ... | ...       | ...           | ...        | ...    | ...             | ...          |
| 4233 | 1    | 50  | 1.0       | 1             | 1.0        | 0.0    | 0               | 1            |
| 4234 | 1    | 51  | 3.0       | 1             | 43.0       | 0.0    | 0               | 0            |
| 4235 | 0    | 48  | 2.0       | 1             | 20.0       | NaN    | 0               | 0            |
| 4236 | 0    | 44  | 1.0       | 1             | 15.0       | 0.0    | 0               | 0            |
| 4237 | 0    | 52  | 2.0       | 0             | 0.0        | 0.0    | 0               | 0            |

4238 rows × 9 columns



In [10]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4238 entries, 0 to 4237
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                   4238 non-null   int64
1   age                    4238 non-null   int64
2   education              4133 non-null   float64
3   currentSmoker          4238 non-null   int64
4   cigsPerDay             4209 non-null   float64
5   BPMeds                 4185 non-null   float64
6   prevalentStroke        4238 non-null   int64
7   prevalentHyp           4238 non-null   int64
8   diabetes               4238 non-null   int64
9   totChol                4188 non-null   float64
10  sysBP                  4238 non-null   float64
11  diaBP                  4238 non-null   float64
12  BMI                    4219 non-null   float64
13  heartRate              4237 non-null   float64
14  glucose                3850 non-null   float64
15  TenYearCHD             4238 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 529.9 KB
```

In [11]: df1=df.fillna(value=0)
df1

Out[11]:

|      | male | age | education | currentSmoker | cigsPerDay | BPMeds | prevalentStroke | prevalentHyp |
|------|------|-----|-----------|---------------|------------|--------|-----------------|--------------|
| 0    | 1    | 39  | 4.0       | 0             | 0.0        | 0.0    | 0               | 0            |
| 1    | 0    | 46  | 2.0       | 0             | 0.0        | 0.0    | 0               | 0            |
| 2    | 1    | 48  | 1.0       | 1             | 20.0       | 0.0    | 0               | 0            |
| 3    | 0    | 61  | 3.0       | 1             | 30.0       | 0.0    | 0               | 1            |
| 4    | 0    | 46  | 3.0       | 1             | 23.0       | 0.0    | 0               | 0            |
| ...  | ...  | ... | ...       | ...           | ...        | ...    | ...             | ...          |
| 4233 | 1    | 50  | 1.0       | 1             | 1.0        | 0.0    | 0               | 1            |
| 4234 | 1    | 51  | 3.0       | 1             | 43.0       | 0.0    | 0               | 0            |
| 4235 | 0    | 48  | 2.0       | 1             | 20.0       | 0.0    | 0               | 0            |
| 4236 | 0    | 44  | 1.0       | 1             | 15.0       | 0.0    | 0               | 0            |
| 4237 | 0    | 52  | 2.0       | 0             | 0.0        | 0.0    | 0               | 0            |

4238 rows × 16 columns



```
In [27]: feature_matrix=df1.iloc[:,0:15]
        target_vector=df1.iloc[:, -1]
```

```
In [28]: feature_matrix.shape
```

```
Out[28]: (4238, 15)
```

```
In [29]: target_vector.shape
```

```
Out[29]: (4238,)
```

```
In [30]: from sklearn.preprocessing import StandardScaler
```

```
In [31]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [32]: logr =LogisticRegression()
        logr.fit(fs,target_vector)
```

```
Out[32]: LogisticRegression()
```

```
In [35]: observation=[[1.4,2.3,5.0,11,12,13,14,15,3,4,5,7,6,7,13]]
```

```
In [36]: prediction=logr.predict(observation)
        print(prediction)
```

```
[1]
```

```
In [37]: logr.classes_
```

```
Out[37]: array([0, 1], dtype=int64)
```

```
In [38]: logr.predict_proba(observation)[0][0]
```

```
Out[38]: 0.00016936367260200758
```

```
In [39]: logr.predict_proba(observation)[0][1]
```


```
Out[39]: 0.999830636327398
```

```
In [40]: import re
        from sklearn.datasets import load_digits
        import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        import seaborn as sns
        from sklearn.linear_model import LogisticRegression
        from sklearn.model_selection import train_test_split
```

```
In [41]: digits = load_digits()
digits
```

```
Out[41]: {'data': array([[ 0.,  0.,  5., ...,  0.,  0.,  0.],
                          [ 0.,  0.,  0., ..., 10.,  0.,  0.],
                          [ 0.,  0.,  0., ..., 16.,  9.,  0.],
                          ...,
                          [ 0.,  0.,  1., ...,  6.,  0.,  0.],
                          [ 0.,  0.,  2., ..., 12.,  0.,  0.],
                          [ 0.,  0., 10., ..., 12.,  1.,  0.])),
          'target': array([0, 1, 2, ..., 8, 9, 8]),
          'frame': None,
          'feature_names': ['pixel_0_0',
                            'pixel_0_1',
                            'pixel_0_2',
                            'pixel_0_3',
                            'pixel_0_4',
                            'pixel_0_5',
                            'pixel_0_6',
                            'pixel_0_7',
                            'pixel_1_0',
                            'pixel_1_1',
                            ...,
                            'pixel_16_7']
          }
```

```
In [42]: plt.figure(figsize=(20,4))
for index,(image,label)in enumerate(zip(digits.data[0:5],digits.target[0:5])):
    plt.subplot(1,5,index+1)
    plt.imshow(np.reshape(image,(8,8)))
    plt.title('Number:%i\n' %label,fontsize=15)
```



```
In [45]: x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,test_s
```

```
In [46]: print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(1257, 64)
(540, 64)
(1257,)
(540,)
```

```
In [47]: logre=LogisticRegression(max_iter=10000) # if error comes declare max_iter=10000
logre.fit(x_train,y_train)
```

```
Out[47]: LogisticRegression(max_iter=10000)
```

```
In [49]: print(logre.predict(x_test))
```

```
[3 6 2 5 4 3 7 6 7 0 5 3 3 2 5 1 1 0 4 5 2 5 5 4 2 2 4 6 8 6 6 8 9 8 6 5 5
 5 2 8 5 6 6 8 7 6 2 2 2 3 4 5 9 5 0 9 1 4 2 6 4 0 3 9 7 2 6 1 9 7 5 7 7 5
 3 2 4 8 7 5 7 6 1 1 5 3 5 2 7 1 8 3 1 1 3 5 7 7 9 2 7 6 2 4 1 5 0 3 7 8 3
 9 4 2 3 4 9 2 6 0 9 6 4 6 5 0 5 4 4 6 5 6 1 3 9 2 6 2 5 8 6 9 5 3 3 5 7 1
 9 1 7 7 7 4 3 1 9 0 0 1 6 1 1 8 4 0 5 9 3 1 3 0 6 7 8 2 2 0 2 3 5 6 4 2 6
 6 2 8 4 7 7 4 8 5 2 6 8 4 0 8 3 9 1 0 2 9 4 2 8 0 9 9 5 1 9 4 4 9 4 7 1 3
 4 6 9 6 3 2 0 9 8 0 7 7 8 1 1 0 6 0 4 5 6 0 0 8 9 8 6 9 3 2 5 8 7 5 6 3 2
 6 6 2 5 7 5 1 9 7 8 4 4 5 7 3 2 8 6 1 2 5 3 8 3 7 0 5 5 3 5 4 3 6 9 5 5 7
 9 5 7 4 2 6 3 9 6 1 8 9 6 2 7 9 4 4 2 2 3 1 5 5 0 2 3 5 8 3 8 3 6 1 9 2 3
 9 7 6 4 8 4 7 2 3 2 2 2 8 5 0 8 7 3 5 7 2 0 9 0 4 6 4 8 3 9 5 8 5 0 7 5 8
 4 6 7 0 0 7 8 8 5 1 2 9 2 9 9 0 3 5 5 5 1 9 3 4 8 9 5 0 0 2 5 6 9 9 7 3 5
 1 9 0 8 3 1 4 0 8 2 0 1 4 0 5 2 8 2 9 8 7 0 2 9 3 8 0 0 2 0 4 2 7 1 1 5 9
 4 2 8 6 4 5 3 6 6 4 3 9 4 2 8 0 8 1 0 4 6 5 7 7 7 3 9 4 8 0 7 5 9 2 7 9 9
 0 0 2 1 7 5 4 9 5 8 0 2 9 1 4 5 8 0 4 2 9 8 5 9 8 9 7 4 6 3 6 9 2 2 2 3 5
 8 8 1 9 3 0 3 7 3 7 2 8 0 0 1 5 1 1 3 3 1 0]
```

```
In [48]: print(logre.score(x_test,y_test))
```

```
0.9685185185185186
```

```
In [ ]:
```